

---

# **KinCat**

***Release 1.2***

**Kyungjoo Kim, Craig Daniels, Habib Najm**

**Oct 13, 2025**



GETTING STARTED:

<b>1</b>	<b>Contents</b>	<b>1</b>
1.1	Introduction . . . . .	1
1.2	Building KinCat . . . . .	1
1.3	Running KinCat . . . . .	6
1.4	File Format . . . . .	9
1.5	KinCatPy . . . . .	19
1.6	KinCat Overview . . . . .	22
1.7	API . . . . .	24



## CONTENTS

### 1.1 Introduction

**KinCat** is an open source software library for parallel kinetic Monte Carlo (KMC) simulations focusing on heterogeneous catalysis applications. In particular, we implement lattice KMC and use Kokkos for its performance portability layer.

---

### 1.2 Building KinCat

For convenience, we explain how to build KinCat using the following environment variables that users can modify according to their working environments. Note that cmake requires use of a separate build directory protecting the source code.

```
# repositories
export KINCAT_REPOSITORY_PATH=/where/you/clone/kincat/git/repo
export KOKKOS_REPOSITORY_PATH=/where/you/clone/kokkos/git/repo
export GTEST_REPOSITORY_PATH=/where/you/clone/gtest/git/repo

# build directories
export KINCAT_BUILD_PATH=/where/you/build/kincat
export KOKKOS_BUILD_PATH=/where/you/build/kokkos
export GTEST_BUILD_PATH=/where/you/build/gtest

# install directories
export KINCAT_INSTALL_PATH=/where/you/install/kincat
export KOKKOS_INSTALL_PATH=/where/you/install/kokkos
export GTEST_INSTALL_PATH=/where/you/install/gtest
```

### 1.2.1 System Requirements

- CMake version 3.16 and higher.
- Compiler supporting C++14 and higher standards.
- Boost library 1.75 and higher.
- Jupyter notebook (or lab) and Python3 for visualization and post-processing.
- sphinx for documentation.
- OpenMPI for running KinCat with multiple data.

#### Mac OSX

We use macports for the standard software distributions of required tools and libraries. Install or check the following tools are installed using macports.

```
sudo port install cmake clang-13 boost python310 py38-sphinx py38-sphinx_rtd_theme_
↪openmpi-clang13

which mpirun-openmpi-clang13
mpirun-openmpi-clang13 is /opt/local/bin/mpirun-openmpi-clang13

which clang++-mp-13
clang++-mp-13 is /opt/local/bin/clang++-mp-13

which cmake
cmake is /opt/local/bin/cmake

export MYCXX=clang++-mp-13

pip install jupyterlab
```

#### Weaver

Weaver was a Sandia National Laboratories compute resource equipped with IBM Power9 processors accelerated by NVIDIA V100 GPU. Although it was recently taken offline, we include notes relating to Weaver, hoping that they may be a useful guide for building and running on other resources. We used the following modules.

```
module purge
module load gcc/11.3.0
module load cuda/11.8.0
module load openmpi
module load cmake
module load openblas
module load ninja
module load git
module load boost
module load hdf5

export MYCXX=${KOKKOS_INSTALL_PATH}/bin/nvcc_wrapper
export INCLUDE=${OPENMPI_ROOT}/include:${INCLUDE}
```

(continues on next page)

(continued from previous page)

```

which gcc
/projects/ppc64le-pwr9-rhel8/compilers/gcc/11.3.0/gcc/8.3.1/base/tchbki3/bin/gcc

echo ${BOOST_ROOT}
/projects/ppc64le-pwr9-rhel8/tpls/boost/1.80.0/gcc/11.3.0/base/wrmyc3o

which cmake
/projects/ppc64le-pwr9-rhel8/utilities/cmake/3.29.6/gcc/8.5.0/base/7ne4ua7/bin/cmake

```

## 1.2.2 How to Build

First, clone Kokkos, GTEST and KinCat code repositories.

```

git clone https://github.com/kokkos/kokkos.git ${KOKKOS_REPOSITORY_PATH}
git clone https://github.com/google/googletest.git ${GTEST_REPOSITORY_PATH}
git clone https://github.com/sandialabs/KinCat.git ${KINCAT_REPOSITORY_PATH}

```

### Kokkos

This builds Kokkos on Intel Haswell architectures and installs Kokkos to `${KOKKOS_INSTALL_PATH}`. For more details, see [Kokkos github pages](https://github.com/kokkos/kokkos). We can use this script for OSX.

```

cd ${KOKKOS_BUILD_PATH}
cmake \
  -D CMAKE_INSTALL_PREFIX="${KOKKOS_INSTALL_PATH}" \
  -D CMAKE_CXX_COMPILER="${MYCXX}" \
  -D Kokkos_ENABLE_SERIAL=ON \
  -D Kokkos_ENABLE_OPENMP=ON \
  -D Kokkos_ENABLE_DEPRECATED_CODE=OFF \
  -D Kokkos_ARCH_HSW=ON \
  ${KOKKOS_REPOSITORY_PATH}
make -j install

```

On Weaver, we compile Kokkos for NVIDIA GPUs. Note that we use Kokkos `nvcc_wrapper` as its compiler instead of directly using the `nvcc` compiler. The architecture flag indicates that the host architecture is IBM Power9 and the GPU architecture is Volta70 generation.

```

cd ${KOKKOS_BUILD_PATH}
cmake \
  -D CMAKE_INSTALL_PREFIX="${KOKKOS_INSTALL_PATH}" \
  -D CMAKE_CXX_COMPILER="${KOKKOS_REPOSITORY_PATH}/bin/nvcc_wrapper" \
  -D Kokkos_ENABLE_SERIAL=ON \
  -D Kokkos_ENABLE_OPENMP=ON \
  -D Kokkos_ENABLE_CUDA:BOOL=ON \
  -D Kokkos_ENABLE_CUDA_UVM:BOOL=OFF \
  -D Kokkos_ENABLE_CUDA_LAMBDA:BOOL=ON \
  -D Kokkos_ENABLE_DEPRECATED_CODE=OFF \
  -D Kokkos_ARCH_VOLTA70=ON \
  -D Kokkos_ARCH_POWER9=ON \

```

(continues on next page)

(continued from previous page)

```
    ${KOKKOS_REPOSITORY_PATH}
make -j install
```

## GTEST

We use GTEST as our testing infrastructure. With the following cmake script, the GTEST can be compiled and installed.

```
cd ${GTEST_BUILD_PATH}
cmake \
  -D CMAKE_INSTALL_PREFIX="${GTEST_INSTALL_PATH}" \
  -D CMAKE_CXX_COMPILER="${MYCXX}" \
  ${GTEST_REPOSITORY_PATH}
make -j install
```

## Boost

The Boost library may be installed by the following script.

```
export BOOST_INSTALL_PATH=/where/you/install/boost
mkdir -p ${BOOST_INSTALL_PATH}
cd ${BOOST_INSTALL_PATH}
curl -L https://boostorg.jfrog.io/artifactory/main/release/1.75.0/source/boost_1_75_0.
tar.bz2 -o boost_1_75_0.tar.bz2
tar -xvf boost_1_75_0.tar.bz2
export BOOST_ROOT=${BOOST_INSTALL_PATH}/boost_1_75_0
```

## KinCat

Build KinCat and link with Kokkos and Gtest. The following script shows how to compile KinCat on Weaver while linking with TPLs explained above.

```
cd ${KINCAT_BUILD_PATH}
cmake \
  -D CMAKE_INSTALL_PREFIX=${KINCAT_INSTALL_PATH}\
  -D CMAKE_CXX_COMPILER="${MYCXX}" \
  -D CMAKE_CXX_FLAGS="-g -I${OPENMPI_ROOT}/include " \
  -D CMAKE_EXE_LINKER_FLAGS="" \
  -D CMAKE_BUILD_TYPE=RELEASE \
  -D KINCAT_SITE_TYPE="char" \
  -D KINCAT_ENABLE_DEBUG=OFF \
  -D KINCAT_ENABLE_VERBOSE=ON \
  -D KINCAT_ENABLE_TEST=ON \
  -D KINCAT_ENABLE_EXAMPLE=ON \
  -D KOKKOS_INSTALL_PATH="${KOKKOS_INSTALL_PATH}" \
  -D GTEST_INSTALL_PATH="${GTEST_INSTALL_PATH}" \
  -D HDF5_INCLUDE_DIRS="${HDF5_INC}" \
  -D HDF5_LIBRARY_DIRS="${HDF5_LIB}" \
  -D HDF5_LIBRARIES="hdf5" \
  ${KINCAT_REPOSITORY_PATH}/src
```



To install KinCat on OSX, we use the following script. Note that in both scripts the HDF5 related keys are only needed if HDF5 functionality is desired.

```
cmake \
-D CMAKE_INSTALL_PREFIX=${KINCAT_INSTALL_PATH} \
-D CMAKE_C_COMPILER="clang-mp-12" \
-D CMAKE_CXX_COMPILER="clang++-mp-12" \
-D CMAKE_CXX_FLAGS="-g" \
-D CMAKE_EXE_LINKER_FLAGS="" \
-D CMAKE_BUILD_TYPE=RELEASE \
-D KINCAT_ENABLE_DEBUG=OFF \
-D KINCAT_ENABLE_VERBOSE=ON \
-D KINCAT_ENABLE_TEST=ON \
-D KINCAT_ENABLE_EXAMPLE=ON \
-D KOKKOS_INSTALL_PATH="${HOME}/kokkos/kokkos_install/release" \
-D GTEST_INSTALL_PATH="/opt/local" \
-D HDF5_INCLUDE_DIRS="/usr/local/hdf5/include" \
-D HDF5_LIBRARY_DIRS="/usr/local/hdf5/lib" \
-D HDF5_LIBRARIES="hdf5" \
${KINCAT_SRC_PATH}
```

A successful installation creates the following directory structure in `${KINCAT_INSTALL_PATH}`. Note that the `site_type` is determined at compile time. In the above cmake configuration, the type is set `char` and the max number of species in KinCat is 256. For a bigger simulation, users can set this `short` or `int`, though this is unlikely to be needed. Also note that currently the `char` setting is overwritten to use `short` instead. This is due to requiring a signed type since KinCatPy uses -1 to indicate an unspecified site in a configuration.

```
- bin
- kincat.x: an executable for solving single problem
- kincat-batch.x: an executable for solving multiple problem with batch parallelism
- plot-dump.ipynb: visualization for jupyter notebook
- examples
- Ru02-dictionary.json : auto-generated from KinCatPy
- Ru02-rates.json : used to set process rates for simulation
- kincatpy
-   - readme.txt
-   - ruo2_input.txt
- non-batch
-   - readme.txt
-   - ex1-input.json
-   - ex2-input.json
-   - plot-ex1.py
-   - plot-ex2.py
- batch
-   - readme.txt
-   - ex3-input.json
-   - plot-ex3.py
-   - input-rates-override-Ru02.json
- include
-   - kincat
-     - header files
- lib (or lib64)
-   - cmake: cmake environment when other software interface KinCat via cmake
-   - libkincat.a
```

(continues on next page)

(continued from previous page)

```
- unit-test
- kincat-test.x: unit test executable
- test-files: sample files that will be used in test
```

## Spack

As an alternative to the ‘manual’ build instructions above, we have also created Spack build instructions. Spack is a package manager intended for scientific software, allowing for software and all of its dependencies to be built with minimal user interaction. Note that this functionality has not been tested extensively. If Spack is not already available, it needs to be set up:

```
git clone --depth=100 --branch=releases/v0.21 https://github.com/spack/spack.git ~/spack
cd ~/spack/
. share/spack/setup-env.sh
```

While Spack often has recognized software packages, for now KinCat can only be built with a local spack package file.

```
spack repo create kincat-local-repo
spack repo add {PATH_TO_SPACK_LOCAL}/kincat-local-repo
mkdir kincat-local-repo/packages/kincat
cp kincat/spack/kincat-package.py kincat-local-repo/packages/kincat/package.py
```

This build file needs to be modified depending on the version of KinCat desired. Both the url reference and the version number and sha256 key should be updated. The version of Kokkos to be used may also be modified. Once Spack and the local repo are set up, KinCat should be able to be installed with just a single command:

Because all the dependencies are also being installed, this may take several minutes. Note that the executable may be saved in an obscure directory.

## 1.3 Running KinCat

We can run the `kincat.x` executable in the `${KINCAT_INSTALL_PATH}/core` directory. The following command line options are available for the current version. There are three verbose options to show additional run-time information. The `verbose-parse` shows raw input from the input json file. Note that configurations and events specification can be different in the dictionary and within the code as KinCat requires a sorted list of configurations. The `verbose` option shows the object details used in KinCat. The `verbose-iterate` option shows the KMC algorithm details. Each verbosity type has four levels: 0) show nothing, 1) basic information, 2) reserved for users verbose comments, 3) lengthy details, 4) show everything.

```
cd ${KINCAT_INSTALL_PATH}/bin
./kincat.x --help
Usage: ./kincat.x [options]
options:
--echo-command-line    bool      Echo the command-line but continue as normal
--help                 bool      Print this help message
--input                string     Input dictionary file name
                              (default: --input=input.json)
--verbose              int       Verbosity level
                              (default: --verbose=0)
```

(continues on next page)

(continued from previous page)

<code>--verbose-iterate</code>	<code>int</code>	Verbosity level <code>in</code> KMC iteration (default: <code>--verbose-iterate=0</code> )
<code>--verbose-parse</code>	<code>int</code>	Verbosity level <code>in</code> parsing step (not yet <code>sorted</code> ) (default: <code>--verbose-parse=0</code> )

Description:  
KinCat Main

The source files include several examples. We show output for example 1 here. By default, we use Kokkos OpenMP as the execution space for CPU processors while this example uses the serial residential time algorithm. Thus, users may want to set `export OMP_NUM_THREADS=1` not to carry some Kokkos overhead of using multiple threads. However, this is unlikely to be significant for this limited run.

```
cd ${KINCAT_INSTALL_PATH}/bin
../../../../../kincat_build/core/kincat.x --input='ex1-input.json' --verbose-iterate=1
Kokkos::OpenMP::initialize WARNING: OMP_PROC_BIND environment variable not set
In general, for best performance with OpenMP 4.0 or better set OMP_PROC_BIND=spread and OMP_PLACES=threads
For best performance with OpenMP 3.1 set OMP_PROC_BIND=true
For unit testing set OMP_PROC_BIND=false
-- Kokkos::OpenMP is used
solver_type = serial-rta
-- Lattice : Lattice
-- number of species : 3
-- domain : [10, 10, 2]
-- edge vectors :
[6.43,0]
[0,3.12]
-- basis :
[0,0]
[0.5,0]

Reset constraints (due to sets construction): 22, 2
instance: 0, 0, [0, 0]
instance: 1, 5, [0, 0]
instance: 2, 1, [0, 0]
instance: 3, 6, [0, 0]
instance: 4, 2, [0, 0]
instance: 5, 14, [-1, 0]
instance: 6, 10, [-1, 0]
instance: 7, 7, [0, 0]
instance: 8, 18, [-1, 0]
instance: 9, 3, [0, 0]
instance: 10, 15, [-1, 0]
instance: 11, 11, [-1, 0]
instance: 12, 8, [0, 0]
instance: 13, 19, [-1, 0]
instance: 14, 4, [0, 0]
instance: 15, 16, [0, 0]
instance: 16, 12, [0, 0]
instance: 17, 17, [0, 0]
instance: 18, 9, [0, 0]
```

(continues on next page)

(continued from previous page)

```

instance: 19, 20, [0, 0]
instance: 20, 13, [0, 0]
instance: 21, 21, [0, 0]
n_process_types found : 22
-- ProcessDictionary : Dictionary
-- variant ordering : (2, 20)
-- configuration list : (27, 4)
-- process instance list : (22, 3), constraints: (22, 2), rates : (22)
-- ProcessDictionary : Instance details
-- process instance list : (22), process list : (22)
-- Solver : serial-rta
-- Dump : Dump
-- filename dump : "dump-ex1.json"
-- filename restart: "restart_sites.json"
creating json file : dump-ex1.json
-- Stats : Stats
-- filename stats : "stats-ex1.json"
epoch = 0, t = 5.00028e-06
-- # of events occurred : 16889
epoch = 1, t = 1.00014e-05
-- # of events occurred : 33363
...

```

This KMC runs and produces the output files `dump-ex1.json` and `stats-ex1.json`. We can post-process these by using `python plot-ex1.py`.

### 1.3.1 Solving for Multiple Samples

For small and medium problem sizes, it may be beneficial to use a batch parallel version of `kincat` i.e., `kincat-batch.x`. Examples 4 and 5 demonstrate the use of this batch version, and more details of the batch parallelism use case will be explained later with the batch input file.

```

cd ${KINCAT_INSTALL_PATH}/bin
./kincat-batch.x --input="ex4-input.json"

```

### 1.3.2 Running on Weaver

To run the code with a GPU, we first allocate an interactive compute node with following command. A single node is dedicated for the user. The Power9 CPU has 40 cores and can utilize 160 threads with symmetric multi-processing (SMP4) accelerated with 4 GPUs. Since Kokkos does not support the multi GPU use case, a user explicitly maps the MPI processes to different GPUs by adding `--kokkos-num-devices=4`.

```

[weaver11] bsub -gpu num=1 -Is -q rhel8 bash
***Forced exclusive execution
Job <42355> is submitted to queue <rhel7W>.
<<Waiting for dispatch ...>>
<<Starting on weaver1>>
[weaver1 ~]$ ./kincat.x --input='input.json'

```

Note that this will use the default GPU, but there are four GPU's available. This can be seen through the following environment variables:

By setting these environment variables to only include a single value, e.g. '1', then that specific GPU will be assigned the job, allowing the user to use all GPU's without interference from concurrent jobs.

The same bsub command, except without the '-gpu num=1' arguments, is used to request a CPU node on Weaver. If desired, the number of threads used can be set by setting the environment variable OMP\_NUM\_THREADS. Note that the KinCat build is specific to whether CPU or GPU is intended to be used. Two builds are required if the user wishes to be able to use both.

### 1.3.3 Restart Simulation

When a simulation completes before it reaches its steady state, the simulation can be restarted using `restart-sites.json` and `dump-batch-sites.json` output. These files contain the last snapshot of the simulation state and they are created at the end of the simulation or when the code catches an exception. See the lattice section of the input file format explained next. These filenames are currently hard-coded into KinCat. They are only produced if a dump-style output is enabled, and are produced with a json format even if HDF5 outputs are requested.

## 1.4 File Format

KinCat requires three input files specifying 1) solver, 2) dictionary, and 3) rates. The solver input can be specified by a user to control the KMC simulations while the dictionary input is auto-generated from a pre-processor, KinCatPy. The rates input is separated as we may want to use/test different set of rates that can be computed by multiple sources e.g., RMG.

### 1.4.1 Main Input

The following json input is the main input file extracted from the RuO2 example 1.

```
{
  "kincat": {
    "dictionary" : "${KINCAT_INSTALL_PATH}/src/examples/RuO2-dictionary.json",
    "rates" : "${KINCAT_INSTALL_PATH}/src/examples/RuO2-rates.json",
    "sites" : {
      "lattice" : [ 10, 10, 2 ],
      "type" : "random",
      "random seed" : 12345,
      "random fill ratio" : [0.0, 0.0],
      "filename" : "dump-sites.json"
    },
    "dump" : {
      "dump filename" : "dump-ex1.json",
      "dump interval" : 2.5E-5
    },
    "statistics" : {
      "stats filename" : "stats-ex1.json",
      "types" : ["species_coverage", "process_counts"],
      "stats interval" : 2E-5
    },
    "solver" : {
      "type" : "serial-rta",
```

(continues on next page)

(continued from previous page)

```

        "random seed" : 13542,
        "random pool size" : 10000,
        "max number of kmc kernel launches" : 1000,
        "max number of kmc steps per kernel launch" : 50000,
        "time range" : [ 0, 1E-3, 5E-6 ]
    }
}
}

```

- kincat object includes the following items:
  - dictionary specifies the file path to the dictionary json file.
  - rates specifies the file path to the rate json file.
  - sites object includes lattice information and its initialization method.
    - \* lattice: 2D lattice object. The array input implies [ x-length, y-length, n-basis ]. The lengths are specified in unit cells, and must be integers.
    - \* type: lattice initialization method.
      - random value fills the lattice randomly with percentages specified by random fill ratio key-word; otherwise, random fill ratio is ignored. random fill ratio should be an array with as many entries as the number of species (excluding vacant sites), and in the same order as they were defined in KinCatPy. The sum of these fill ratios should not exceed one. A random seed is given and used just for the initialization routine.
      - file value loads site specification from a file given by the restart file filename; otherwise, filename is ignored. Note that while the initial configuration may be read-in from a file, the initial simulation time is always set by time begin as set below.
  - dump object relates to an output file with the full simulation state.
    - \* dump filename object specifies the name of the dump output file.
    - \* dump interval object indicates the minimum time between outputs of the full simulation state.
  - statistics object relates to an output file of various desired statistics for the KMC simulation.
    - \* filename object specifies the name of the statistics output file.
    - \* types object is a list of the various statistics desired by the user. Currently implemented types include "process\_counts", "species\_coverage", and "site\_species\_coverage". The "process\_counts" option outputs the number of times each process has occurred up to that point in the KMC simulation which can be used to calculate turn-over frequencies. The "species\_coverage" option outputs the fraction of the lattice sites filled by each species, and "site\_species\_coverage" option further breaks down this fraction by each unique type of lattice site. Further types may be added later, or added by the user.
    - \* stats interval object indicates the minimum time between outputs of the statistics listed in types.
  - solver object include the following items to specify the KMC simulation.
    - \* type specifies the algorithm selected: serial-rta, sublattice, batch-rta.
      - serial-rta solver implements the BKL, n-fold, or 1st Order KMC algorithm.
      - batch-rta solver implements the serial-rta algorithm over multiple samples.
      - sublattice solver is based on the synchronous sublattice algorithm of Shim and Amar.

- \* `domain` specifies the lattice size [`x-length`, `y-length`] processed by a single process (or a group of threads) in a parallel algorithm (`sublattice`). If a parallel algorithm is not used, `domain` is ignored.
- \* `random seed` specifies the random seed used in the KMC simulations.
- \* KinCat uses a random number generator and pre-generates an array of random numbers and `random pool size` indicates the array size. The minimum random number array is  $2 * (\# \text{ of subdomains}) * (\text{max number of KMC iterations per kernel launch})$ .
- \* The simulation will complete either after it meets `max number of kmc kernel launches` or the simulation reaches `time end`.
- \* Each KMC kernel launch will complete either after the `max number of kmc steps per kernel launch` or the `dt` is reached. Note that if a kernel completes before reaching `dt`, the limiting `dt` will not be updated and the next kernel will only proceed until the prior `dt` is reached. This helps maintain regular output intervals and explains why ‘extra’ kernels may be desirable.
- \* `time range` specifies the [`time begin`, `time end`, `dt (time-increment)`].

For both the dump and statistics objects, giving a filename of “none” will result in no output file of that type. Filenames must either use “.json” or “.hdf5” extensions. If KinCat is built with HDF5 functionality, it will detect which file extension is used and output that file format.

Note that a single `solver.advance()` function runs until it reaches `dt` time step or the maximum number of KMC iterations per kernel launch. When `dt` is set zero, the code runs for the specified number of KMC steps. If a user wants to ensure the `dt` time step is reached for each kernel launch, then the number of KMC steps needs to be sufficiently large. If the number of steps is not large enough to reach the desired `dt` timestep, then another kernel will launch with the same limiting endtime as the prior kernel. For the `serial-rta`, running the code setting without the `dt` constraint does not cause any simulation errors.

The sublattice algorithm as implemented requires that the lattice size be an even integer multiple of sublattice domain size in both x and y dimensions. Also, the sublattice domain size should be larger than 2X the interaction range of the lattice model (included in the dictionary file) to avoid potential corruptions. The sublattice algorithm requires synchronization among subdomains. A subset of subdomains are evolved simultaneously while the others are frozen. The algorithm rotates through subsets until all subdomains are synchronized at `dt`, rejecting the final steps that would extend past `dt`. The timestep `dt` should be sufficiently small so that the kinetics within the subdomain will not be significantly different at the end than at the beginning. Otherwise it will lead to significant errors. On the other hand, shorter timesteps lead to more numerous rejected events and reduced efficiency. KinCat will output a warning if a significant number of events occur before timestep `dt`. However, this warning level is somewhat arbitrary. The user is strongly encouraged to carefully consider what parameters will balance errors and efficiency for their system.

## 1.4.2 Ensemble Input

To exploit `kincat-batch.x`, it is required to append an ‘ensemble’ section to the above main input. We use the input of Example 4 to illustrate it.

```
"ensemble" : {
  "number of samples" : 4,
  "solver random number variations" : {
    "apply" : "enabled",
  },
  "sites random variations" : {
    "apply" : "enabled",
    "random fill ratio" : {
      "0" : [0.3, 0.5],
      "1" : [0.4, 0.4],
    }
  }
}
```

(continues on next page)

(continued from previous page)

```

        "2" : [0.5, 0.3],
        "3" : [0.6, 0.2]
    },
    "rates variations" : {
        "apply" : "enabled",
        "type" : "file",
        "processes" : ["CO_ads_cus", "CO_ads_br"],
        "process rates" : {
            "0" : [1.85e+06 , 2.15e+06 ],
            "1" : [1.90e+06 , 2.10e+06 ],
            "2" : [1.95e+06 , 2.05e+06 ],
            "3" : [2.00e+06 , 2.00e+06 ]
        },
        "process instances" : [0 , 1],
        "instance rates" : {
            "0" : [1.85e+06 , 2.15e+06 ],
            "1" : [1.90e+06 , 2.10e+06 ],
            "2" : [1.95e+06 , 2.05e+06 ],
            "3" : [2.00e+06 , 2.00e+06 ]
        },
        "override filename" : "../input-rates-override-RuO2.json"
    }
}

```

- ensemble object includes the following items:
  - number of samples specifies the number of concurrent simulations to be run.
  - When solver random number variations is enabled, each sample uses a different sequence of random numbers to select KMC events.
  - site random variations varies the initial configurations of samples according to the random fill ratios.
    - \* random fill ratio contains dictionary objects used to define the species fill ratios as in the sites object, but with each sample uniquely defined. If site random variations is enabled but the random fill ratio object is not present, each sample will initialize with a different initial configuration, but with the fill ratio set included in the ‘sites’ object.
    - \* When this option is disabled, samples use the same initial configuration, which is specified in the sites object.
  - process rates variations allows for samples to use different process rates when it is enabled.
    - \* type can be either inlined or file.
      - inlined looks for processes and process instances, either or both or which may be included.
      - file type value take override filename keyword to load the user specified rates for samples. The format of the file should be a json file with the processes and/or process instances arrays and their corresponding rates dictionaries.
    - \* processes keyword indicates an array of the processes which rates are to be modified. If the processes keyword is present, then the process rates must also be present.
    - \* The process rates object is a dictionary with the sample index and an array of rates. The rates correspond to the processes listed in the processes array. There must be the same number of rates



provided for each sample as processes listed.

- \* `process instances` keyword indicates an array of the process instances which rates are to be modified. If the `process instances` keyword is present, then the `instance rates` must also be present.
- \* The `instance rates` object is a dictionary with the sample index and an array of rates. The rates correspond to the process instances listed in the `process instances` array. There must be the same number of rates provided for each sample as processes instances listed.

Note that the `batch-rta` solver type is the only solver that supports the ensemble section for now. The batch input should be used with `kincat-batch.x` executable. Also note that if a configuration file is provided in the `sites` object, the samples will be initialized from there rather than any fill ratios provided in the `ensemble` object. This configuration file may include either the same number of samples as the current batch, or may include only one, in which case all samples will be initialized with the same configuration.

### 1.4.3 Dictionary Input

This file contains the information needed to set up the lattice shape and define possible KMC events. It is generated as an output of KinCatPy and should not need to be modified by the user. This is an explanation of the information contained in the dictionary file. KinCatPy can be used to generate a dictionary file for three use-cases. KinCat uses so-called 1) ‘Sets’, 2) ‘Uniconfig’, and 3) full symmetry (‘Fullsym’) configurations specifying the event mechanism. The styles give statistically equivalent results, but use differing reliance on symmetry operations which changes how the processes and configurations are catalogued and retrieved. The Fullsym case specifies each possible process in a single configuration, while the sets and uniconfig cases use the symmetry of the lattice to reduce the size of the configuration(s) catalogued. The following example dictionary file specifies the sets case, which divides the processes by the sites they involve and define multiple configurations.

```
{
  "crystal": {
    "edge vectors": [[6.43, 0.0], [0.0, 3.12]],
    "basis vectors": [[0.0, 0.0], [0.5, 0.0]],
    "symmetry operations": {
      "shape": [4, 6],
      "data": [1, 0, 0, 1, 0.0, 0.0, -1, 0, 0, -1, 0.0, 0.0, -1, 0, 0, 1, 0.0, 0.0,
↪ 1, 0, 0, -1, 0.0, 0.0]
    },
    },
    "configurations": {
      "interaction range": [2, 2],
      "site coordinates": [[0.0, 0.0], [0.5, 0.0], [0.0, 1.0], [0.5, 1.0]],
      "variant orderings": [[0, 1, 2, 3], [2, 3, 0, 1]],
      "symmetry orderings": {"0": [[0], [0]], "1": [[0], [0]], "2": [[0, 1], [1, 0]], .
↪ ..},
      "configuration_sets": [0, 3, 6, 12, 18, 27],
      "shape": [27, 4],
      "data": [-1, 0, -1, -1, -1, 1, -1, -1, -1, 2, -1, ...],
      "process dictionary": {
        "processes": ["CO_ads_cus", "CO_ads_br", "O_ads_cus_cus", "O_ads_br_br", ...],
        "process constraints": [[[1, 0, 2]], [[0, 0, 2]], [[1, 0, 1], [3, 0, 1]], [[0, 0,
↪ 1], ...],
        "process symmetries": [[0], [0], [0], [0], [0, 1], [0], [0], [0], [0], [0, 1], ↪
↪ [0, 1], ...],
        "shape": [22, 3],
        "data": [0, 2, 0, 2, 0, 5, 3, 5, 1, 5, 3, 6, 6, 9, 2, 7, 7, 14, 8, 8, ...]
      }
    }
  }
}
```

(continues on next page)

(continued from previous page)

```
}
}
```

- **crystal** object describes a unit cell structure and its symmetry operations.
  - **edge\_vectors** includes two vectors that form a unit cell (parallelogram) that is repeated to tile the lattice domain.
  - **basis\_vectors** represents the position of sites in the unit cell coordinates.
  - **symmetry\_operations** provides rotation matrices and translation vectors that produce equivalent symmetry configurations.
    - \* **shape** indicates [ # of symmetries, array size (4 entries for 2x2 matrix, 2 entries for 2x1 vector) ] and is used to interpret the data array.
- **configurations** object includes a list of possible configurations. A configuration is defined as a unique arrangement of simulation species within the interaction range of a system process.
  - **interaction\_range** defines the range around the central site that needs to be recalculated after each event due to possible changes to processes and rates in that region. It is given in unit cells. It also affects the minimum domain size for parallel solvers.
  - **site\_coordinates** includes the position of sites in the reference configuration. The position is given in units of the crystal edge vectors.
  - **variant\_orderings** and **symmetry\_orderings** relate to the possible enumerations (or re-ordering of sites) forming symmetry equivalent configurations when the configuration is mapped to the lattice sites.
  - **configuration\_sets** stores the divisions between different configuration definitions in the catalogue.
  - The list of configurations is stored as a 2D array [ # of configurations, configuration size ] where the entries of data are the species index. The sets dictionary style may result in some site occupations being specified as '-1'. This indicates that this site is not important for that configuration definition.
- **process\_dictionary** object describes process mapping from one configuration to the other configuration.
  - **processes** is a list of process labels. For the reduced symmetry configurations, the processes are unique.
  - **process\_constraints** is a rank-3 array [ # of processes, # of constraints, constraint size(3) ]. For a corresponding process, a constraint [ site index, initial species, final species ] represents the initial and final species of the specified site. The specified site in the initial and final configurations should match those given in the constraint. Otherwise, we do not consider it as a valid process instance. The initial and final species may be the same, indicating that the site is not changed by the process, but that it is a 'bystander' site important for the process definition but not changed by the process.
  - **process\_symmetries** indicates a pattern index that should be accounted when computing valid events. For example, an absorption event can be counted multiple times in the reduced symmetry configurations. To prevent this multiple counting, the process symmetry information include [0] pattern index so that only the first symmetry pattern is used when searching for possible events. Without the process symmetry information, the KMC process will find multiple events that are equivalent (for the adsorption example, it would find four events: one for each symmetry operation). To avoid the duplicated event search, we can also use the Fullsym dictionary style as explained in the next section.
  - A process instance is specified as [initial configuration, final configuration, process index] and stored as rank-2 array [# of events, event size(3)].

A Fullsym input is shown below. Note that the event dictionary grows exponentially with the number of sites and the number of species. Using the Fullsym or even Uniconfig dictionary style might be prohibitive for a large reaction model.

```

{
  "crystal": {
    "edge vectors": [[6.43, 0.0], [0.0, 3.12]],
    "basis vectors": [[0.0, 0.0], [0.5, 0.0]],
    "symmetry operations": {
      "shape": [1, 6],
      "data": [1, 0, 0, 1, 0.0, 0.0]
    }
  },
  "configurations": {
    "interaction range": [2, 2],
    "site coordinates": [[-0.5, 0.0], [0.0, -1.0], [0.5, -1.0], [0.0, 0.0], [0.5, 0.0], ↵
    ↵ [0.0, 1.0], [0.5, 1.0], [1.0, 0.0]],
    "variant orderings": [[0, 1, 2, 3, 4, 5, 6, 7]],
    "symmetry orderings": {"0": [[0, 1, 2, 3, 4, 5, 6, 7]]},
    "configuration_sets": [0, 6561],
    "shape": [6561, 8],
    "data": [0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, ...]
  },
  "process dictionary": {
    "processes": ["CO_ads_cus", "CO_ads_br", "O_ads_cus_cus", "O_ads_br_br", ...],
    "process constraints": [[[4, 0, 2]], [[3, 0, 2]], [[4, 0, 1], [6, 0, 1]], ...],
    "process symmetries": [[0], [0], [0], [0], [0], [0], [0], [0], [0], [0], ...],
    "shape": [32076, 3],
    "data": [0, 54, 0, 0, 162, 1, 0, 30, 2, 0, 90, 3, 0, 108, 4, 0, 2268, 5, 1, 55, 0, 1, ..
    ↵ . ]]
  }
}

```

Here, we explain the major differences from the Sets case.

- The Fullsym style, which does not use symmetries, has an identity matrix for `symmetry operations`.
- The `variant ordering` is trivial i.e., identity map.
- The `configuration_sets` is trivial.
- As expected, the number of possible configurations and process instances is much bigger than the Sets case e.g., 6561 vs 27 and 32076 vs 22.
- `events` can have duplicated processes names e.g., same absorption process with different initial configurations.
- `process constraints` and `process symmetries` require inputs for the same number of `processes` array size.
- `process symmetries` is trivial and the Fullsym case does not have multiple event searches.

## 1.4.4 Rate Input

The rate input is explained with the sample script below.

```

{
  "default rate": 0,
  "process specific rates" : {
    "CO_ads_cus" : 2.04E+06 ,
    "CO_ads_br" : 2.04E+06,

```

(continues on next page)

(continued from previous page)

```

    "O_ads_cus_cus" : 3.81E+03,
    "O_ads_br_br" : 3.81E+03,
    "O_ads_br_cus" : 3.81E+03,
    "CO_des_cus" : 1.82E+07,
    "CO_des_br" : 5.50E+04,
    ...
  },
  "event specific rates": {
    "5" : 1.85E+07,
    "11" : 2.00E+06,
    ...
  }
}

```

- **default rate** is used when the rate is not otherwise specified. A negative default value can be used for error checking if the input file must specify all processes (or process instances).
- **process specific rates includes rates for processes.**
  - All process instances with the same process will be set to the same rate.
  - Events with processes not specified will be set to the default rate.
- **process instance specific rates includes rates for specific instances.**
  - Process instances not specified will be set to the process rate if one was specified, or the default rate otherwise.

### 1.4.5 Dump Output

When **dump** is enabled from the main input, the code dumps the output of the sites in the following format.

```

{
  "samples" : 4,
  "number of species": 3,
  "coordinates": {
    "shape": [ 800, 2 ],
    "data": [ 0, 0, 3.215, 0, 0, 3.12, 3.215, 3.12, 0, 6.24, ... ]
  },
  "sites": [
    {
      "sample": 0,
      "time": 0,
      "data": [ 2, 1, 1, 2, 1, 1, 2, 2, 2, 2, 1, 1, 2, 1, ... ]
    },
    {
      "sample": 1,
      "time": 0,
      "data": [ 1, 2, 2, 0, 2, 1, 0, 2, 1, 2, 0, 1, 2, 2, ... ]
    },
    {
      "sample": 2,
      "time": 0,
      "data": [ 1, 2, 1, 1, 0, 2, 1, 2, 2, 2, 1, 1, 0, 1, ... ]
    }
  ]
}

```

(continues on next page)

(continued from previous page)

```

    },
    {
      "sample": 3,
      "time": 0,
      "data": [ 1, 1, 1, 1, 1, 1, 0, 0, 1, 2, 1, 1, 1, 0, ...]
    },
    {
      "sample": 0,
      "time": 5.00036e-07,
      "data": [ 2, 2, 1, 2, 1, 2, 2, 2, 2, 2, 1, 2, 2, 0, ...]
    },
    {
      "sample": 1,
      "time": 5.00061e-07,
      "data": [ 1, 2, 2, 2, 2, 1, 2, 2, 1, 2, 2, 0, 2, 2, ...]
    },
    {
      "sample": 2,
      "time": 5.00135e-07,
      "data": [ 1, 2, 1, 2, 2, 2, 1, 2, 2, 2, 1, 2, 2, 2, ...]
    },
    {
      "sample": 3,
      "time": 5.00028e-07,
      "data": [ 1, 1, 1, 2, 1, 2, 2, 0, 1, 2, 1, 1, 1, 2, ...]
    },
    ...
  ]
}

```

A dump file can be used for post-processing and it includes 1) number of species, 2) coordinates, and 3) time series of sites information. The time series information includes the sample number (even if not running an ensemble), the time stamp, and the current occupation state at each site. Similar information is given with hdf5 style outputs. An example of post-processing is given with each example. Additionally, `restart-sites.json` and `dump-batch-sites.json` files are created to record the last site configurations when the code completes, which can be used for restarting the simulation.

### 1.4.6 Stats Output

When `statistics` is enabled from the main input, the code dumps the selected statistics in the following format.

```

{
  "lattice size" : [ 20, 20, 2 ],
  "samples" : 4,
  "number of species" : 3,
  "number of processes" : 22,
  "processes" : [ "CO_ads_cus", "CO_ads_br", "O_ads_cus_cus", "O_ads_br_br", ... ],
  "readings" : [
    {
      "sample" : 0,
      "time" : 0,

```

(continues on next page)

(continued from previous page)

```

    "species coverage" : [ 0.2, 0.3, 0.5 ],
    "process counts" : [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0 ],
    ↪ 0, 0 ],
    "site species coverage" : [[ 0.22, 0.29, 0.49 ], [ 0.18, 0.31, 0.51 ]]
  },
  {
    "sample" : 1,
    "time" : 0,
    "species coverage" : [ 0.2, 0.4, 0.4 ],
    "process counts" : [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0 ],
    ↪ 0, 0 ],
    "site species coverage" : [[ 0.22, 0.4175, 0.3625 ], [ 0.18, 0.3825, 0.4375 ]]
  },
  {
    "sample" : 2,
    "time" : 0,
    "species coverage" : [ 0.2, 0.5, 0.3 ],
    "process counts" : [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0 ],
    ↪ 0, 0 ],
    "site species coverage" : [[ 0.22, 0.48, 0.3 ], [ 0.18, 0.52, 0.3 ]]
  },
  {
    "sample" : 3,
    "time" : 0,
    "species coverage" : [ 0.2, 0.6, 0.2 ],
    "process counts" : [ 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0 ],
    ↪ 0, 0 ],
    "site species coverage" : [[ 0.19, 0.625, 0.185 ], [ 0.21, 0.575, 0.215 ]]
  },
  {
    "sample" : 0,
    "time" : 4.00134e-07,
    "species coverage" : [ 0.0425, 0.21125, 0.74625 ],
    "process counts" : [ 2352, 160, 1, 2, 6, 2222, 4, 0, 0, 0, 0, 103, 0, 0, 0, 4, 0,
    ↪ 0, 11, 0, 0, 78 ],
    "site species coverage" : [[ 0, 0.315, 0.685 ], [ 0.085, 0.1075, 0.8075 ]]
  },
  {
    "sample" : 1,
    "time" : 4.00335e-07,
    "species coverage" : [ 0.04125, 0.2925, 0.66625 ],
    "process counts" : [ 2153, 166, 3, 1, 4, 2001, 3, 0, 0, 0, 0, 94, 0, 0, 0, 8, 0,
    ↪ 0, 19, 0, 0, 83 ],
    "site species coverage" : [[ 0.005, 0.4325, 0.5625 ], [ 0.0775, 0.1525, 0.77 ]]
  },
  {
    "sample" : 2,
    "time" : 4.00002e-07,
    "species coverage" : [ 0.0425, 0.37375, 0.58375 ],
    "process counts" : [ 1876, 196, 1, 0, 9, 1716, 8, 0, 0, 0, 0, 77, 0, 0, 0, 9, 0,
    ↪ 0, 11, 0, 1, 109 ],
    "site species coverage" : [[ 0.0025, 0.5, 0.4975 ], [ 0.0825, 0.2475, 0.67 ]]
  }

```

(continues on next page)

(continued from previous page)

```

    },
    {
      "sample" : 3,
      "time" : 1.60017e-06,
      "species coverage" : [ 0.03375, 0.395, 0.57125 ],
      "process counts" : [ 8205, 202, 7, 1, 8, 7901, 13, 0, 0, 0, 0, 61, 0, 0, 0, 36,
→0, 0, 73, 0, 0, 123 ],
      "site species coverage" : [[ 0, 0.65, 0.35 ], [ 0.0675, 0.14, 0.7925 ]]
    },
    ...
  ]
}

```

In the `readings` object, the `sample` and `time` will always be present. However, the rest of the objects will depend on the selected options. The `species coverage` gives current fill fraction of each species over all sites. The `site species coverage` breaks this fill fraction down by lattice site. The `process counts` gives the count of each process since the start of the simulation, which can be used to calculate turn-over frequencies.

## 1.5 KinCatPy

**KinCatPy** is an open source Python script to write dictionary input files for use with KinCat. It receives inputs describing the lattice, possible species and process mechanisms, and the range of possible interactions. It outputs a KinCat dictionary file which contains information about all possible symmetrically unique process instances and all configurations needed to define those (given the input parameters). A given process mechanism, such as desorption from the surface, may have different rates due to lateral interactions with species on nearby sites. KinCatPy defines a unique process instance for each process that may occur between two unique initial and final configurations. Thus, rates for each instance can be adjusted to account for the lateral interactions, giving the user as much specificity in setting rates as desired. KinCatPy is adapted from KineCluE, another open source code intended to calculate transport coefficients of mobile clusters. Note that while KineCluE was written to handle bulk crystals, KinCatPy is only intended for 2D lattice descriptions.

KinCatPy is run in the command line window or terminal as `python ./path_to_kincatpy/kincatpy.py input.txt`.

### 1.5.1 Main Input

Some of the input language and structure remains the same as KinCluE (<https://github.com/lukamessina/kineclue>). While we outline the inputs here, the user may find KineCluE documentation to be helpful background. It is also helpful to understand how the Bravais lattice construction allows any lattice to be defined by a repeating unit cell and basis vectors within that cell. The following input is extracted from the RuO2 example.

```

& CRYSTAL test # creation of the crystal with 2 vectors of 2 components
6.43 0.0
0.0 3.12
& BASIS s 2
0 0 0
1 0.5 0
& UNIQUEPOS 2 # describes the sites in the crystal that the atoms and defect will occupy.
s 0 0

```

(continues on next page)

(continued from previous page)

```

s 0.5 0
& RANGE 3.0 # the float is the interaction radius
& SPECIES 2
2 1 1 0 0.1
2 1 1 CO 0.1
& PROCMECH
#adsorption processes
%% 1 CO_ads_cus
s 0.5 0 0 > 2
%% 1 CO_ads_br
s 0 0 0 > 2
...
#desorption processes
%% 1 CO_des_cus
s 0.5 0 2 > 0
%% 1 CO_des_br
s 0 0 2 > 0
...
##diffusion processes
%% 2 CO_cus_cus
s 0.5 0 2 > 0
s 0.5 1 0 > 2
%% 2 CO_br_br
s 0 0 2 > 0
s 0 1 0 > 2
...
##Recombination/desorption of CO2
%% 2 CO_cus_O_cus
s 0.5 1 1 > 0
s 0.5 0 2 > 0
%% 2 CO_br_O_br
s 0 1 1 > 0
s 0 0 2 > 0
...

```

It is useful to recognize two possible coordinate systems that KinCatPy will accept. The first is the standard x-y or Cartesian coordinates. The second is in terms of lattice edge vectors which define the unit cell that tiles the plane. These are referred to as orthogonal and supercell coordinate systems and ‘o’ and ‘s’ respectively are used to specify which is used in the input file. Also recognize that & is used to indicate a new input command and # may be used to comment the remainder of the line.

- & CRYSTAL is used to define the unit cell of the lattice. The string after the command can be used to name the lattice if the user wishes. The next lines comprise two edge vectors that define the lattice unit cell. These values are always given in the orthogonal coordinate system, and these values are used to map between the orthogonal and supercell coordinate systems moving forward.
- & BASIS command is used to define the basis vectors of the lattice. If it is not present, then a single basis vector of (0,0) is assumed. The coordinate system of the basis vectors (‘o’ or ‘s’) must be specified, and the number of basis vectors given. For each basis vector, a line of the form `site_type a b` is required. The `site_type` is an integer, and can be used to indicate if the site is fundamentally the same or different than another which is important for symmetry considerations. The `a b` terms are the coordinates of the basis vector in the specified coordinate system.
- & UNIQUEPOS command is similar to the & BASIS command, and must always be included. The BASIS vectors



are used to define the symmetry of the system, while UNIQUEPOS vectors define sites the species can occupy. For UNIQUEPOS, only one vector for each symmetrically equivalent site is required. Thus, there should be the same number of vectors here as there are unique site types given in the BASIS section. In UNIQUEPOS, the number of vectors is again required, but the coordinate system is defined for each vector, in the form  $s(o) \ a \ b$ .

- & RANGE command is used to define the interaction range, and the distance is in the units of the orthogonal coordinate system. Each configuration will include all sites within this range of a process, so increasing it may dramatically increase the complexity of the KMC model generated.
- & SPECIES command is used to give information about the possible species in the system. Currently, each species is considered as a single point, so it is impossible to include information about orientation or multi-site binding in the model. The number of species needs to be given, and then each species is specified by a line of the form `n pos_bools... name size`. The number of each species is intended to be provided by `n` (as used in KineCluE). However, the user just needs to ensure that this integer is larger than the number of that species specified in any single JUMPMECH. Next, `pos_bools` refers to a set of 0/1 flags indicating if that species can occupy the UNIQUEPOS given. There should be as many flags as there are UNIQUEPOS specified. The `name` is a unique text string. The last number is `size` indicating the implied radius of the species (orthogonal units). Species with sufficient size may 'block' neighboring sites and prevent other species from occupying them. The size may be set arbitrarily small if this functionality is not desired.
- & PROCMECH command is used to define possible process mechanisms through sets of constraints. After the & PROCMECH command is given, a new process mechanism definition is denoted by a line of form `% n_constraints name`. Each process name should be unique. `n_constraints` is the number of constraints that define the process. Each constraint on the process is given by a line of the form `s(o) \ a \ b \ ini_species_type > fin_species_type`. The letter  $s(o)$  is used to denote which coordinate system will be used for `a \ b`, which are the site coordinates. `ini_species_type` and `fin_species_type` are the integer indices of the initial and final species respectively. The species inputted using the & SPECIES command are assigned indices in their list order, beginning with one. The species index 0 is reserved to indicate a vacant site. Note that each constraint definition consists of a site and the change of species on that site. Thus, defining a diffusion jump involves two constraints, one for the species disappearing from the initial site and another for the species appearing in the final site. So called 'bystander' or 'spectator' species that do not change may also be included in the constraints. For example, a reaction between two 'A' species may be catalyzed by the presence of a 'B' species. The B species would need to be included in the constraints with an appropriate coordinate relative to the A species, but the initial and final `species_type` of that constraint would be the same. Note that a only single symmetry of a process needs to be included. Symmetric processes are found by accounting for the symmetries of the lattice.

& DIRECTORY command can be used to specify an output directory where the output files will be generated. If this command is not included, an output directory named 'CALC/' is auto-generated in the working directory.

& FULLSYM command can be used to create a dictionary where the configurations and process instance dictionaries are not reduced by symmetry operations.

& UNICONFIG command can be used to create a dictionary where the configurations and process instance dictionaries are reduced by symmetry to a single configuration template. Symmetric processes are recovered by symmetry operations.

If neither the FULLSYM or UNICONFIG flags are included, KinCatPy will use the default 'Sets' dictionary style. In this style, processes are grouped by involved sites to create multiple configuration templates. Symmetric processes are again recovered by symmetry, and this dictionary style will produce the smallest catalogues.

Note again that the complexity of the model given will greatly depend on 1) the number of species included, 2) the range of interactions accounted for (set by RANGE), and 3) the choice of process definitions. The computational demands of KinCatPy are largely determined by the number of unique configurations that need to be specified. Since it will identify all possible permutations of the species arranged on a set of lattice sites, the number of species and number of sites in the configuration definition will both have an exponential impact on the number of configurations overall. The configuration includes all lattice sites within the interaction range of a site that is changed by a process (the species on that site changes, not just that it is included in a bystander constraint). Increasing the interaction range will increase the number of sites needed for the configuration definition. However, for the reduced symmetry cases, careless process definition may

also increase the number of sites defined. The process definitions should use the symmetries that overlap as much as possible. In this example defining CO adsorption to ‘cus’ site process using site 0 0.5 and the CO desorption from ‘cus’ site process using site 0 -0.5 is valid, since the sites are symmetrically equivalent to each other. However, when compared to defining both processes with the same site, this would lead to an expansion of the configuration definition and increase the complexity of the model defined by KinCatPy. Since all symmetries of each process are included in the full symmetry case, it does not matter. We recommend using the sets case for most simulations, but the full-symmetry case may be easier for the user to interpret.

We also note that for large model dictionaries, a multi-thread script (kincat\_event\_calc.py) is called. However, this only significantly improves efficiency with large numbers of configurations, and so is not called if the model has fewer than 100,000 configurations. Even models with over a 100,000 configurations and over 1,000,000 process instances only takes a few minutes to generate. We also note that KinCatPy assumes that the kincat\_event\_calc.py file will be stored in the directory above where KinCatPy is running.

## 1.6 KinCat Overview

A basic KMC algorithm is below. This is often called the resident time algorithm (RTA) and is inherently sequential. In KinCat, we consider three solver implementations: 1) serial reference implementation, 2) sublattice parallel implementation, and 3) batch parallel implementation.

```
SolverSerialRTA::advance(in: t, in: t_step, in: n_max_iterations,
                        in: lattice, in: dictionary, in: rates, in: rates_scan) {
    /// update event rates on lattice
    updateEventRatesLattice(in: lattice, in: dictionary, out: rates);

    for (iter=0; iter<n_max_iteration && t<t_end; ++iter) {
        /// perform prefix sum
        scanRatesLattice(in: rates, out: rates_scan);

        getRandomNumber(out: pi, out: zeta);
        sum_rates = rates_scan(in: rates.extent(0));
        dt = log(zeta)/sum_rates;
        rate_to_search = sum_rates * pi;

        /// cid is the cell index corresponding to a specific rate (sum_rates*pi) in the_
        ↪ scanned array
        searchCellIndex(in: rate_to_search, in: rates_scan, out: cid);

        /// select an event in the selected cell
        findEvent(in: cid, in: dictionary, out: event);

        /// update the cell configuration
        updateLattice(in: cid, in: event, out: lattice);

        /// local updates of rates on the neighborhood of the selected cell
        updateEventRates(in: lattice, in: dictionary, out: rates);

        /// advance time
        t += dt;
    }
}
```

The sublattice algorithm decomposes the lattice into multiple subdomains. In each subdomain, a sequential RTA algorithm runs and time is updated asynchronously. To avoid the potential conflicts on the boundary regions, the algorithm updates a single set of independent domains at a time.

```
SolverSublattice::advance(in: t, in: t_step, in: n_max_iterations,
                          in: lattice, in: dictionary, in: rates, in: rates_scan) {

  /// update event rates on lattice
  updateEventRatesLattice(in: lattice, in: dictionary, out: rates);

  /// even odd selector
  quad = { {0,0}, {1,0}, {0,1}, {1,1} };

  for (iter=0; iter<n_max_iteration && t<t_end; ++iter) {
    for (q=0; q<4; ++q) {
      parallel_for(in: team_policy(n_domains), in: LAMBDA(member) {
        /// compute domain indices corresponding to the member's league rank
        getDomainIndex(in: member, in: lattice, out: d0, out: d1);

        /// select even or odd domain in 2D
        if (d0%2 == quad[q][0] && d1%2 == quad[q][1] && t(d0,d1) < t_step) {
          /// get domain specific containers
          auto rates = rates_lattice(d0, d1);
          auto rates_scan = rates_scan_lattice(d0, d1);

          /// nested parallel scan over rates
          parallelScanRates(in: member, in: rates, out: rates_scan);

          /// update event
          single([=]() {
            getRandomNumber(out: pi, out: zeta);
            sum_rates = rates_scan(in: rates.extent(0));
            dt = log(zeta)/sum_rates;
            rate_to_search = sum_rates * pi;

            /// cid is the cell index corresponding to a specific rate (sum_rates*pi) in
            ↪ the scanned array
            searchCellIndex(in: rate_to_search, in: rates_scan, out: cid);

            /// select an event in the selected cell
            findEvent(in: cid, out: event);

            /// update the cell configuration
            updateLattice(in: cid, in: event, out: lattice);

            /// update the domain clock
            t(d0,d1) += dt;
          });

          /// local updates of rates on the neighborhood of the selected cell
          parallelUpdateEventRates(in: member, in: lattice, in: dictionary, in:
            ↪ interaction_range, out: rates);
        }
      });
    }
  }
}
```

(continues on next page)

(continued from previous page)

```
    });  
  }  
}
```

The batch RTA solver runs the serial RTA, but in multiple independent simulations or samples. This may be desired for statistical or UQ analysis. For the batch mode, we implement a separate executable.

---

## 1.7 API

Here, we manually and briefly document important things only.

### 1.7.1 Kokkos Things

Note that some functions are decorated with `KOKKOS_INLINE_FUNCTION`. This means that the function can be called on device. All captured object in the device function is `const` object. We cannot alter the members nor call non-const functions.

`KOKKOS_LAMBDA` cannot capture the implicit or explicit “this” pointer. If a member variable needs to be used in the lambda function, assign the member variable to a local variable. Most of Kokkos view does soft copy and no actual cost for the assignment.

### 1.7.2 class Lattice

The lattice object provides means for accessing specific location of sites and providing symmetry patterns. These patterns are used to construct an configuration array which is used for comparing against dictionary to identify the configuration and possible events.

- `value_type_2d_view<site_type,device_type> _sites`. - The `site_type` is determined at compile time i.e., `short int`, `int` to save bytes for frequent memory transfer between host and device. - `_sites(sid,cid)` stores the current status of species occupied on the site location (`cid`) for the sample (`sid`).
  - `copySites`. When 2d site view is used, it overwrites for all samples. For 1d site view input with sample `id`, it overwrites sites for the specific sample.
  - `getCellIndex(k0, k1, out: cid)`. For a given lattice index `k0` and `k1`, it computes the corresponding cell index. The cell index is different from site index and site index accounts for multiple basis points. Note that we use `k0` and `k1` indices indicating the location of the lattice and we use `l0` and `l1` indices for local domain (sublattice algorithm).
  - `getDomainCellIndex(k0, k1, out: d0, out: d1, out: lid)`. For a given `k0` and `k1`, it computes domain index `d0` and `d1` and site location (`lid`).
  - `getLatticeCellIndex(d0, d1, l0, l1, out: k0, out: k1)`. The function computes lattice index from domain indices.
  - `getSiteIndex(k0, k1, ib, out: site_index)`. The function returns the `site_index` in 1d array. This function needs basis index (`ib`).
  - `adjustPeriodicBoundary(inout: k0, input: k1)`. The lattice assumes periodic conditions and input `k0` and `k1` are adjusted accordingly.

- `getCoordinates`. Used for post processing.

### 1.7.3 class ProcessDictionary

- `value_type_2d_view<ordinal_type, device_type> _variant_orderings`. This includes correct enumeration orders for configuration.
- `value_type_2d_view<site_type, device_type> _configurations`. `_configurations(conf_id, conf_index) = species` on this configuration and index. The configuration array is stored after sorted for binary search.
- `value_type_2d_view<ordinal_type, device_type> _processints`. `_processints(initial_conf_id, final_conf_id, process_id)`. The process instance is stored after it is sorted for binary search. The instances change the current configuration to the final configuration and its mechanism is explained by the process.
- `value_type_2d_view<ordinal_type, device_type> _constraints`. Certain variant ordering of the configuration is legit and this constraint view filter out non-feasible configuration.
- `value_type_2d_view<real_type, device_type> _rates`. `_rates(sample_id, event_id)` indicates a specific fate for the event of the sample. Each sample can use the same or different rates.
- `searchConfiguration(input_key, out: index_configuration, out: index_variant)`. The function performs binary search and report corresponding configuration index and its variant index.
- `searchFirstEvent(index_configuration, out: index_first_event)`. The function reports the index of the first process instance for the given configuration. When system rate is scanned, each cell has a symmetry unique configuration and this function find a list of possible process instances.

### 1.7.4 class SolverBase

A base class for different solver variants i.e., serial-rta, batch-rta, sublattice.

- `initialize` setup the solver internal objects. Since the lattice and dictionary object uses 2d views that can be used for multiple samples, single problem (non-batch mode) can use the same data structure by setting number of samples is 1.
- `advance(2d_view/1dview, ...)`. The function provides two interface. 2d the time view is used for sublattice algorithm and the serial version just use `t_in(0,0)` only. The 1d time array is for the batch mode `t_in(sample_id)`.
- `createSolver(solver_type)`. The function creates solver object based on the input string (`solver_type`).

### 1.7.5 class Dump

- `initialize` function dump the header (lattice coordinates).
- `finalize` function close the dump file and create the last snapshot of sites which can be used for restarting of the simulations.
- `snapshot` function dump the current site information with sample and time stamp.

### 1.7.6 class Stats

-`initialize` function outputs the header of a statistics json file, relevant to the statistics desired. -`finalize` function closes the statistics file and creates the last snapshot of the statistics. -`snapshot` function outputs the chosen statistics to the statistics file with the sample and time stamp.

### 1.7.7 class ProcessCounter

- `initialize` function create counter views for both device and host. The internal objects will be properly deleted as they are reference counted objects.
  - `reset` function zeros all counters.
  - `syncToHost` perform deep copy to host so that a host function can access updated counter information.
  - `update(sample index, event index)` function will increase the counter for the sample.
-